
Priority Queues and Binary Heaps (Part 1)

Binary heaps, types of heaps, and complexity analysis

Connor Baker

Priority Queues and Binary Heaps (Part 1)

Review: Queues

- First in, first out (FIFO)
- Operations:
 - `getFront()`, `enqueue(T t)`, and `dequeue()`
 - `size()` and `isEmpty()`
- Applications
 - Simulate a process with FIFO ordering
 - Scheduling the queue in a CPU, disk, or printer
 - Buffer for file I/O, network communication, or other transmissions
- A lot of the time, tasks in a queue have priorities
 - Dequeue should remove or return the one with the best priority
- Common priority queue operations
 - `add(T t, int p)` and `enqueue(T t, int p)`: enqueue item `t` with priority `p`
 - `peek()` and `findMin()`: return the object with the best priority
 - * Per convention, lower is better
 - * Symmetric code if higher is better
 - `dequeue()` and `deleteMin()`: remove and return the object with the best priority

Priority Queue Implementation

Data Structure	<code>enqueue(T t)</code>	<code>peek*</code>	<code>dequeue*</code>	Notes
Unsorted List	$O(1)$	$O(n)$	$O(n)$	best priority can be any location
Sorted Array	$O(n)$	$O(1)$	$O(1)$	best priority at high index
Sorted Linked List	$O(n)$	$O(1)$	$O(1)$	best at head or tail
Multiple Queues	$O(1)$	$O(m)$	$O(m)$	-
Binary Search Tree	$O(\text{height})$	$O(\text{height})$	$O(\text{height})$	min at left-most

- *: assuming best priority
- n : the number of items in a queue
- m : the number of priority levels

Binary Heap

- A *binary heap* is a binary tree but *not* a binary search tree
- Differences:
 - Sort of sorted: each node is smaller than, or equal to, both its children
 - Must be a complete binary tree




Figure 1: BST Example




Figure 2: Binary Heap Example

Sorted Binary Heap Example




Figure 3: Sorted Binary Heap Example

Heap Order

- Max heap
 - The node is always larger than, or equal to, any of its descendants
- Min heap
 - The node is always smaller than, or equal to, any of its descendants
- Idea: we want to find the min, or max, quickly
 - Keep at the root of the tree
 - Recursive definition: every subtree should have the largest, or smallest, item at the root of the subtree

Binary Heap Examples

- Is it a heap or not?
 - If it is, what kind of heap is it?




Figure 4: Examples

Complete Trees

- Could only be missing nodes in their bottom row
- Nodes in the bottom row are as far left as possible




Figure 5: Incomplete Trees




Figure 6: Complete Trees

Trees and Heaps in Arrays




Figure 7: A complete binary tree and its array representation

- The root is at index 1
- Binary tree
 - $\text{left}(i) = 2 * i$
 - $\text{right}(i) = 2 * i + 1$
 - $\text{parent}(i) = i / 2$

Road Map

- Priority queue
 - $\text{insert}(T, t), \text{findMin}(), \text{deleteMin}()$
- Heap
 - Complete binary tree
 - Heap order
 - * Min heap
 - * Max heap
 - Operations and complexity
 - * Insert: percolate up
 - * Delete: percolate down
- Heap sort

Priority Queue Operations with Binary Heaps

- Use an internal `T array[]` for queue contents
 - Maintain min-heap order in array
 - Make sure it is always a complete tree
- `T findMin()`
 - `return array[root()];`
- `insert(T t, int p)`
 - Insert at the end of the array, increment size
 - * Might violate the min-heap order property
 - Fix by swimming the new element up (percolate up)
- `deleteMin()`
 - Simply removing the root will leave a hole
 - We can swap the last value and the root to fill the hole
 - * `null` out the last value (which prevents loitering)
 - * Decrement the size
 - * The new root *might* not be minimal
 - Percolate the new root value down the tree
- Max heap follows the same ideas

Binary Heap Demo

- Starting from a min-heap like Example 2
 - Insert 50, 18, and 10
 - `deleteMin()`
- Starting from an empty max-heap
 - Insert 2, 3, 5, 3, and 9
 - `deleteMax()` five times

Operation Details

- Basic questions
 - With whom do we compare or swap?

- When do we stop moving?
- Percolate up (bubble up)
 - Compare / swap with the parent
 - Halting condition: when we reach the top (the root) or no longer violate the heap order
- Percolate down (sink down)
 - Compare / swap with a child
 - Halting condition: when we reach the bottom (a leaf) or no longer violate the heap order

Weiss Code Example

```
1  /**
2   * Removes the smallest item in the priority queue.
3   * @return the smallest item.
4   * @throws NoSuchElementException if empty.
5   */
6  public T remove() {
7      T minItem = element();
8      // Move the tail element to the root
9      array[1] = array[currentSize--];
10     // Sink the new root down to fix the heap order
11     percolateDown(1);
12 }
13
14 /**
15  * Internal method to percolate down in the heap.
16  * @param hole the index at which to percolate begins.
17  */
18 private void percolateDown(int hole) {
19     int child;
20     T tmp = array[hole];
21     // Decide which child to compare/swap
22     for (; hole * 2 <= currentSize; hole = child) {
23         child = hole * 2;
24         if (child != currentSize && compare(array[child+1], array[child]) < 0) {
25             child++;
26         }
27         // Keep swapping with children until parent-child comparison result is
28         // satisfactory, or the bottom is reached
29         if (compare(array[child], tmp) < 0) {
30             array[hole] = array[child];
31         } else {
```



```

32     break;
33 }
34 array[hole] = tmp;
35 }
36 }

```

Complexity Analysis

- `findMin()` is clearly $O(1)$
- What about `insert(T t)` and `deleteMin()`?
 - Percolation does most of the work
 - Worst case: $O(\text{height})$
 - * Complete binary tree: $O(\log(n))$
- Note: no `get(T t)` or `remove(T t)`

Priority Queues Comparison

Data Structure	<code>enqueue(T t)</code>	<code>peek*</code>	<code>dequeue*</code>	Notes
Unsorted List	$O(1)$	$O(n)$	$O(n)$	best priority can be any location
Sorted Array	$O(n)$	$O(1)$	$O(1)$	best priority at high index
Sorted Linked List	$O(n)$	$O(1)$	$O(1)$	best at head or tail
Multiple Queues	$O(1)$	$O(m)$	$O(m)$	-
Binary Search Tree	$O(\text{height})$	$O(\text{height})$	$O(\text{height})$	min at left-most
Binary Heap	$O(\log(n))$	$O(1)$	$O(\log(n))$	best priority at root

- *: assuming best priority
- n : the number of items in a queue
- m : the number of priority levels