
Red-Black Trees (Part 2)

Red-black trees, more rotation cases, insertion, and complexity

Connor Baker

Red-Black Trees (Part 2)

Review

- Color properties
- Insertion: add new values as a red leaf
 - Black parent
 - Red parent and red uncle
 - * Recoloring
 - Red parent and black uncle
 - * Rotation and recoloring

Rotation General Cases

- The new node is red, the parent is red, and the uncle is black
 - Note: more black on the uncle's side of the tree
- Step 1: perform a rotation at the parent *if needed*
 - **When do we do this? Why?**

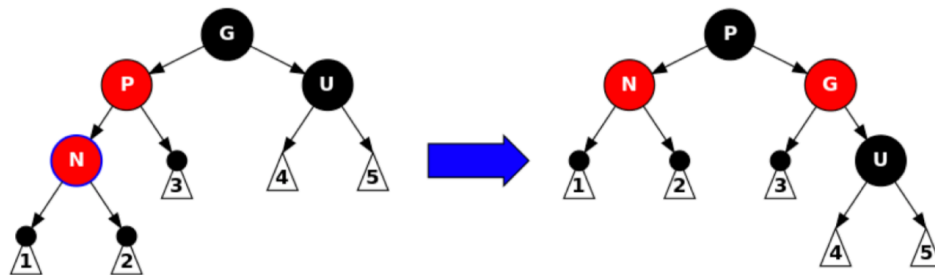


Figure 1: Step 1

- Step 2: perform a rotation at the grandparent
- Step 3: swap the old parent and the grandparent's colors

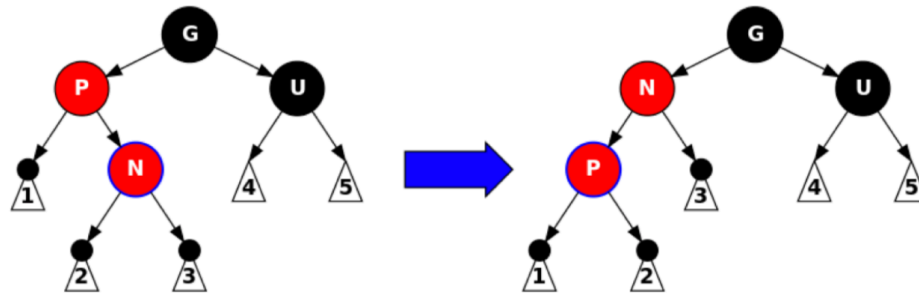


Figure 2: Steps 2 and 3

- We could potentially need an additional step in certain cases
 - Consider the case where we need to use a left-rotation with the newly inserted node

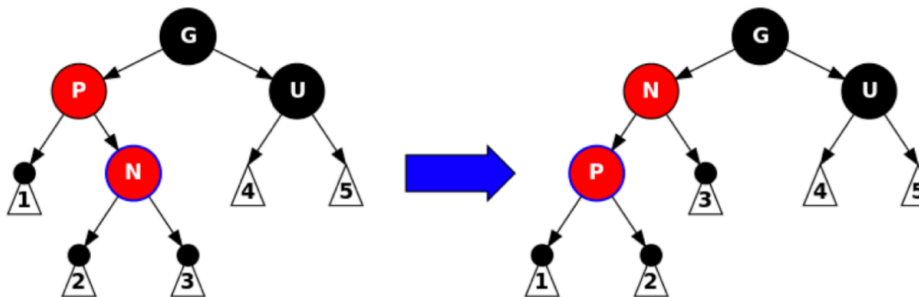


Figure 3: An extra rotation may be needed

Problems with Red Subtree Roots

- Remaining issue: if a recolor makes a subtree root red, then we may have create two consecutive red nodes!
- Strategy one: detect and travel back up to perform additional fixes
 - We can always change the root to black for a final fix
 - We do however have to go up to fix the nodes with rotation or recoloring

Example

- Suppose we want to insert 45 into the following red-black tree.

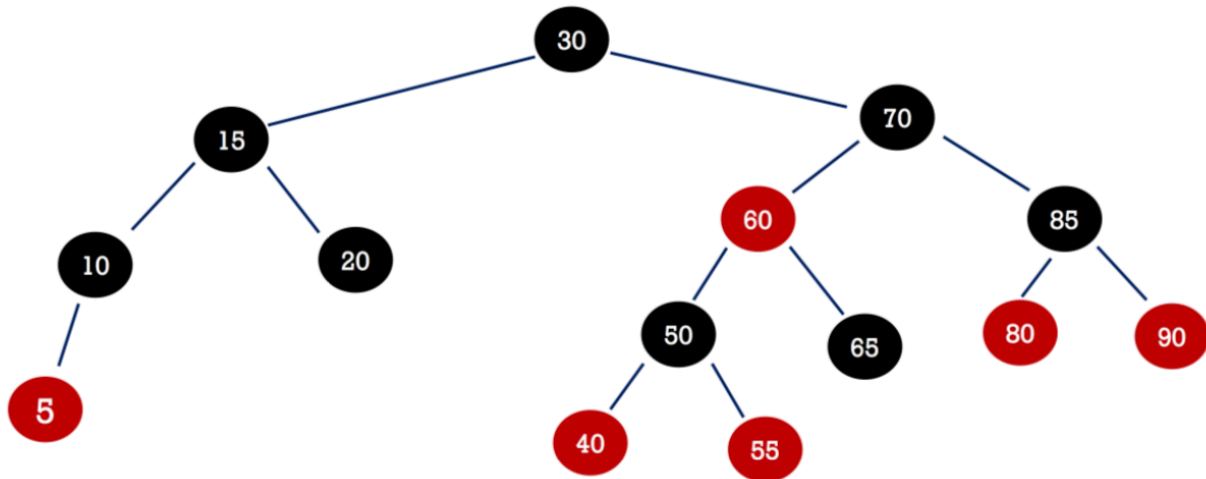
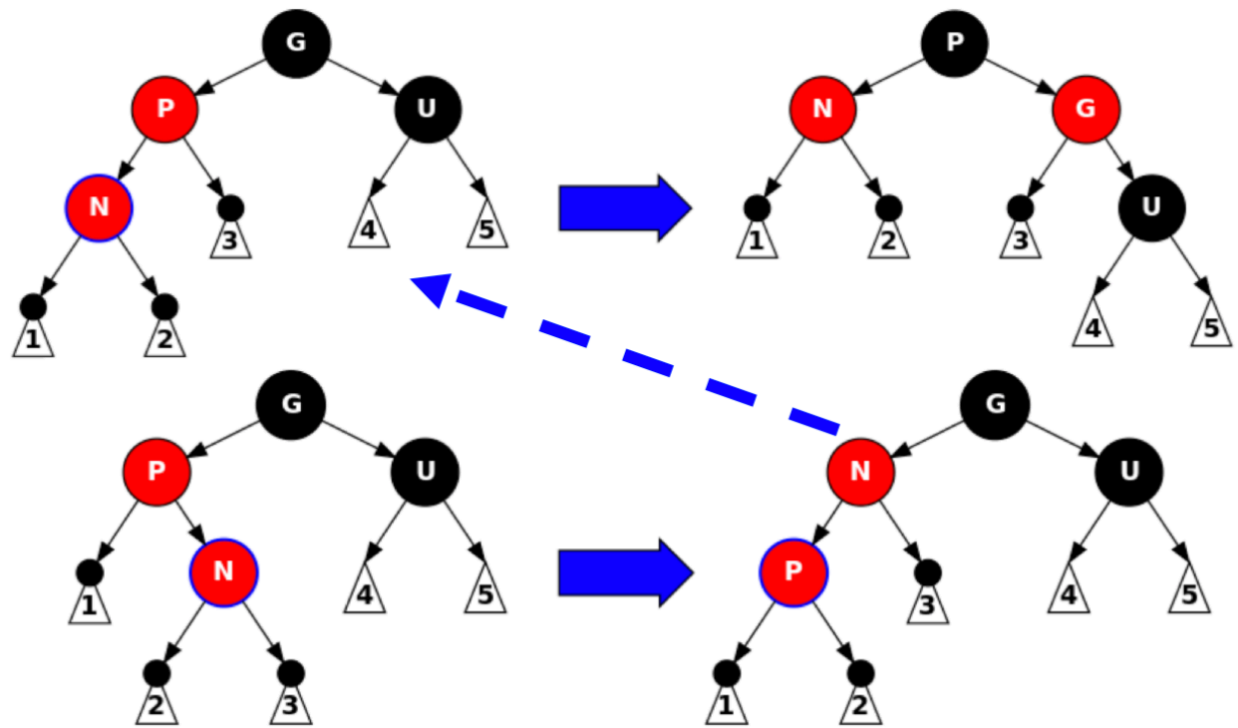
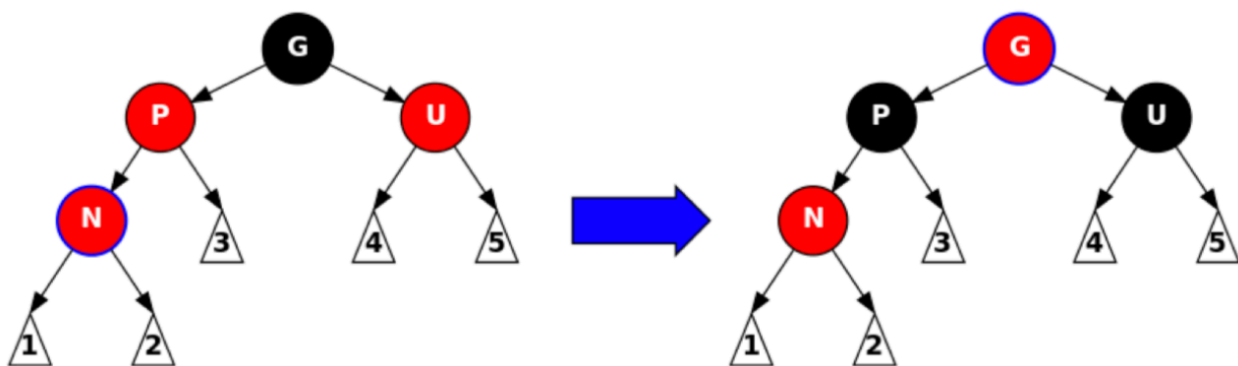


Figure 4: A red-black tree

- Trying to do so would land us with a red-red combo, where the uncle is also red
 - As such, we'd need to recolor them
 - This requires (possibly) more rotations and recoloring up the tree

Top-Down Approach

- Strategy two: down only (top-down insertion)
 - In a single downwards pass, recolor and rotate to ensure that the insertion will succeed without us having to walk back up the tree to fix it afterwards
- What might cause trouble for us so that we need to go back up the tree to fix it?
 - Black parent: easy, no issue there
 - Red parent
 - * **Black uncle?**
 - * **Red uncle?**

Black Uncle**Figure 5:** Cases with a black uncle**Red Uncle****Figure 6:** Case with a red uncle

Top-Down Insertion

- The fix: guarantee that the red parent does not have a red sibling
 - One the way down we need to check the (black) node n
 - If both children are red, change the children to black and change n to red
 - ★ **What if n is the root?**
 - If the parent of n is red, use a single/double rotation and recoloring to fix, then continue down the tree
 - ★ **Can n have a red uncle?**
 - Ensure that after the red insertion, we only need to perform local adjustments
 - ★ There is no percolating back up

Example (Continued)

- Back to our previous example of trying to insert 45 into a red-black tree:

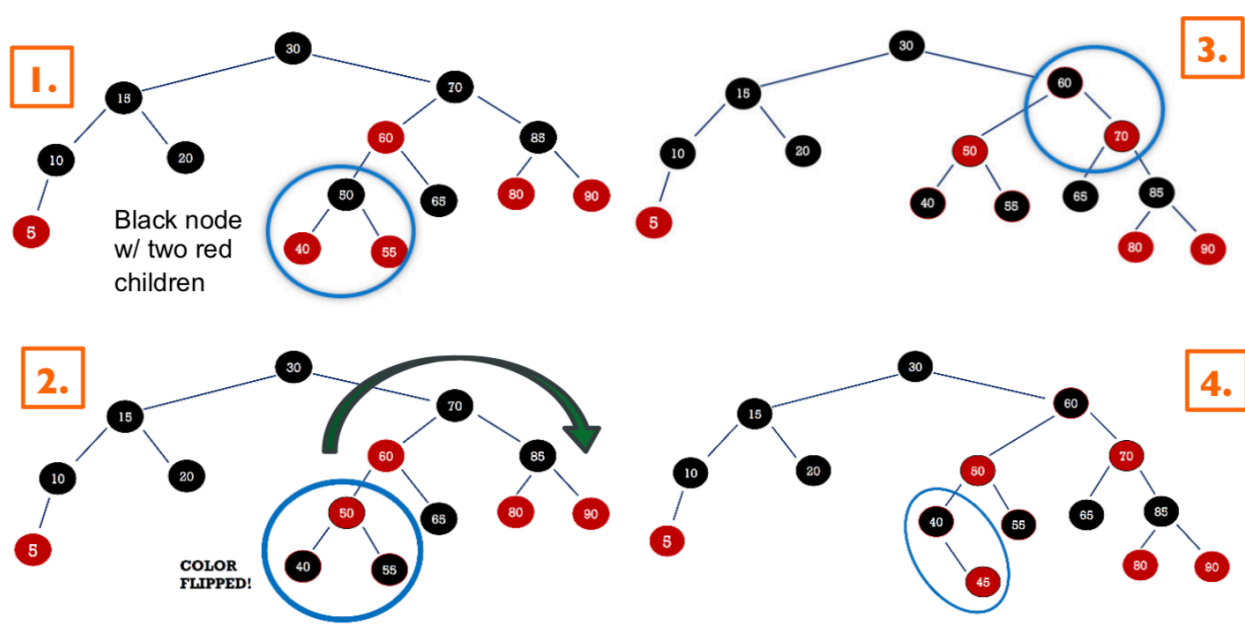


Figure 7: Top-down insertion with a red-black tree

Complexity

- A red-black tree which contains n nodes has a height of $O(\log(n))$

- A detailed proof is available on Wikipedia's Red-Black Tree page
- Notes
- The black height (bh) of a node n counts the number of black nodes from n to a **null** link (not counting n if n is black)
 - Given $bh(n)$, the shortest path from n to any **null** link has $bh(n)$ edges: all of which are black nodes
 - Given $bh(n)$, the longest path from n to any **null** link has $2 * bh(n)$ edges: the nodes alternate between red and black
 - $bh(n) \geq h(n) / 2$, where $h(n)$ is the height of n where $h(\text{null}) = 0$
 - The height of a node n : $h(n)$ is bounded above by $2 * bh(n)$
 - Number of nodes rooted at n : bounded below by $2^{bh(n)} - 1 \leq N$
 - * Proof by induction on height and bh
 - $h(n) \leq 2 \log_2(N + 1)$ where N is the size of the tree

AVL Trees vs. Red-Black Trees

AVL Trees

- Pros
 - Simple(r) to implement
 - Faster lookup (maintains optimal height)
- Cons
 - Slower to insert or delete (because it must maintain optimal height)
 - Simple implementation is recursive and relies on down-up rotations

Red-Black Trees

- Pros
 - Faster insert/delete (doesn't have to maintain optimal height)
 - Implementation tricks can do down-only insertion and deletion
- Cons
 - Complex algorithm
 - Slower lookup (doesn't have optimal height)
 - Implementation is complicated

Balanced BST Benefits

- Keep data in order
- Provided guarantee $O(\log(N))$ find/add/remove
- Reproduce sorted order via an in-order traversal
- Reproduce sorted slices of data
 - Locate a record in $O(\log(N))$ time
 - In-order traversa from that record

Sets and Maps Review

- Closely related data structures
 - A collection of (usually distinct) values
 - Supports efficient look-up
 - * Do we have this value or not?
 - * Is there a value associated with this key>
- Conventions:
 - Do not allow keys to be null
 - Do not allow duplicate values for a key
 - * `put(key1, val1)` followed by `put(key1, val2)` sees `val2` overwrite `val1`
 - Iteration is allowed on the keys
 - * We can the use the key to get the associated value

Maps and Sets Implementations

- Main concern: can we efficiently handle a large number of `get` operations after a large number of `put` or `get` operations?
- Naive implementations
 - Unordered linked list
 - Ordered array
- Trees
 - Binary search trees
 - Balanced binary search trees (AVL, Red-Black, AA)
- Hash tables

Maps and Sets Big- O

Implementation	Worst Case	Worst Case	Average Case	Average Case	Ordered	Remarks
	Search	Insert	Search	Insert		
Unordered List	$O(n)$	$O(n)^1$	$O(n)$	$O(n)^1$	no	
Ordered List	$O(\log(n))$	$O(n)$	$O(\log(n))$	$O(n)$	yes	
HT Chaining	$O(n)$	$O(n)$	$O(n/m)$	$O(n/m)$	no	often used ²
HT Probing	$O(n)$	$O(n)$	$O(1)$	$O(1)$	no	
BSTs	$O(n)$	$O(n)$	$O(\log(n))$	$O(\log(n))$	yes	easy
AVLs	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	yes	easy
Red-Black Trees	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	yes	often used ²

- n : number of values
- m : number of entries
- ¹: $O(n)$ to check for duplicates, $O(1)$ if this is not needed for some reason
- ²: Good constants and relatively easy to implement; used in many libraries

Next Lecture

- Disjoint sets: union/find
- Reading: Chapter 24