
Final Review

A review of all the concepts covered since the midterm

Connor Baker

Final Review

Reading List Notes and Quick References

- Big-O:
 - Chapter 5: Algorithm Analysis
 - Section 5.3: You are not required to prove the theorems, but you need to understand them.
- Lists:
 - Section 6.1-6.3, 6.5: General Lists
 - Chapter 15: Dynamic Array / Array List and Iterator
 - Chapter 17: Linked List
 - Section 6.6: Stack/Queue Concepts
 - Chapter 16.1-16.3: Stack/Queue Implementation
- Trees:
 - Section 18.1: Tree Basics, Binary Trees
 - Section 7.1-7.3: Recursion
 - Section 18.3-18.4: Tree Recursions, Tree Traversals
- Hashing:
 - Section 20.1-20.2: Hashing Basics and Hash Functions
 - Section 20.5: Separate Chaining
 - Section 20.3-20.4: Open Addressing
 - Section 20.6-20.7: Comparisons and Applications
- Graphs:
 - Section 14.1: Graph Basics
 - Section 14.2-14.3: Shortest-path Problem
- Priority Queues / Heap:
 - Section 6.9: Priority Queues
 - Section 21.1-21.3: Binary Heap
 - Section 21.5: Heap Sort
- Binary Search Trees:
 - Section 19.1,19.3: Binary Search Tree Basics
 - Section 19.4: AVL Trees
 - Section 19.5: Red-black Trees
 - Section 19.8: B-Trees
- Disjoint sets:
 - Section 24.1-24.5: Union Find

Sample Questions for Review

NB: Questions before the midterm are the same as they were for the midterm review, which makes sense given that the final is cumulative. Since I published a separate midterm review, these questions are omitted here.

Hashing (Post-Midterm)

45. Compare and contrast the three types of probing we covered for open addressing. Make sure to note the benefits and disadvantages of each.

Well there aren't really *three* different types of probing – rather, there's three ways that we can handle collisions. We have separate chaining and open addressing. Open addressing in turn has two different ways of handling collisions, linear probing and quadratic probing. As a chart, we'd have

1. Separate Chaining
2. Open Addressing
 1. Linear probing
 2. Quadratic probing

Separate Chaining

The benefit of using separate chaining is that collisions just turn into insertions into some data structure (whether it be a binary tree, another `HashMap`, or something else entirely). Generally separate chaining is preferred. However, that means we need to maintain a separate data structure, have another level of indirection, and adding elements to the data structure may require resizing.

Additionally operations like `add(T t)`, `remove(T t)`, and `contains(T t)` have an upper bound of $O(\log(n))$ since we must search whatever data structure we're using for each entry in the hash table. (Using a second `HashMap` keeps the amortized cost at $O(1)$, while using a binary search tree introduces a cost of $O(\log(n))$, for example).

In class we used a linked list as the secondary data structure, which was kind of silly – sure it can grow infinitely and is an easy visual to grasp, but the efficiency just isn't there for the operations that we need (fast `add(T t)`, `remove(T t)`, and `contains(T t)`). For Big- O analysis with separate chaining and linked lists, see my notes from February 28, page 3.

Open Addressing

Linear and quadratic probing search for the next available space in a hash table instead of creating a data structure to handle collisions. As an example, if slot 5 is filled, linear probing will check the sequence $\{5+1, 5+2, 5+3, 5+4, 5+5, \dots\} = \{6, 7, 8, 9, 10, \dots\}$ and so on, while quadratic probing will check $\{5+1, 5+2^2, 5+3^2, 5+4^2, 5+5^2, \dots\} = \{6, 9, 14, 21, 30, \dots\}$.

This introduces a need to use tombstones, markers on each entry that indicate whether it has been removed or not, so that we avoid removing entries which are in search chains.

Linear probing suffers from primary clustering since it continually searches within a small neighborhood of cells. Additionally, it has a load of

$$\frac{1}{2} \left(1 + \frac{1}{(1 - \text{load})^2} \right)$$

which yields a very steep exponential growth curve for the number of cells checked during an insertion, with respect to the load.

Quadratic probing still needs to use tombstones. However, it doesn't suffer from primary clustering since it checks a more dispersed group of cells. The complexity of quadratic probing isn't known.

A major benefit to use quadratic probing is that if the table size is prime and the load factor stays been 0.5, you are guaranteed to be able to insert an item. Additionally, we get the property that no cell is probed twice.

46. Create a hash function for [strings | numbers | cards | people] which generates relatively unique values.
47. Hash [strings | numbers | cards | people] into a hash table of size [some number using [separate chaining | open addressing w/ linear probing | open addressing w/ quadratic probing | open addressing w/ double hashing where $h_2(\text{key}) = 5 - (\text{key} \% 5)$] and the hash function you wrote above.
48. After hashing the above [strings | numbers | cards | people] remove three of them, making sure to visually indicate any special states.
49. Rehash your tables above to [a prime number over double the size].
50. Write code for an [add | remove | contains | rehash | getLoad] method in a hash table which uses [separate chaining | open addressing with linear probing | open addressing with quadratic probing | open addressing with double hashing where $h_2(\text{key}) = 7 - (\text{key} \% 7)$].
51. Explain how Java [String | Integer | Double | Character] class produces hash codes.

Graphs

52. Explain the following terms: graph, node, edge, adjacent to, directed edge, weight, $|V|$, $|E|$, path, simple path, path length, cycle, degree/indegree/outdegree of a vertex, DAG.
53. Explain the difference between the following types of graphs: directed/undirected, weighted/unweighted, cyclic/acyclic, dense/sparse, connected/disconnected.
54. Explain the two ways to store a graph we covered in class. When would you want to use which?
55. Given a graph, draw the corresponding adjacency [matrix | list]. Given an adjacency [matrix | list], draw the corresponding graph.

Graph Algorithms

55. Given a graph and a starting location, perform a [breadth-first | depth-first] traversal showing the steps of the accompanying [queue | stack] data structure. Given a choice of neighbors, choose nodes in numerical order.

56. Write code for a recursive depth-first search. Assume you are given an adjacency matrix `m` (`int[][]`), the `id` of a node to start from, and an `id` of a node to search for. This should return `true/false` depending on if one can get to the search target from the starting node.
57. Given a graph and a source node `id`, perform Dijkstra's shortest path algorithm on a given graph showing the steps. Use a table to track the current status, distance, and parent pointer for each node.

Graph Algorithms (Prof. Russel's section only) – Omitted

Priority Queues/ Heaps

NB: The numbering in the PDF I'm follow is weird – it's not me.

56. What are the common operations supported by a priority queue?
57. Compare and contrast different data structures we covered in class to implement a priority queue.
58. Write the code to [enqueue | dequeue] from a priority queue stored as [a sorted dynamic array | an unsorted linked list | a heap]. Make sure to maintain the optimal Big-O for the given queue storage.
59. How can you determine the [min | max] value in a heap? Show the steps to remove the [min | max] value from the heap. Show the steps to insert a value into a heap and keep the heap order.
60. Given an array representation of binary heap, show the corresponding tree structure. Given a binary heap tree representation, show the corresponding array assuming root is at index [0|1].
61. What is the difference between “delete” in a heap and “delete” in other types of trees we discussed in class?
62. Explain how heapsort is $O(n \log(n))$ and can be done “in place”.
63. Given an array of numbers, show the steps of the optimal heapify algorithm.

All Trees (Binary, K-ary, Binary-Search, AVL, Red-Black, Heap, B/B+)

64. Draw a valid [tree we covered in class].
65. Given a tree (as either a picture or an array), determine if it is a valid [tree we covered in class] and, if not, determine what rule is violated and where the error is.
66. Explain the “rules” of a [tree we covered in class]. What properties have to be maintained when [adding a node to | removing a node from] the tree?
67. Given a [tree we covered in class] and a value, show the steps of [searching for | inserting | deleting] that value.
68. Compare and contrast the [search | insertion | deletion] times for each of the trees we covered.
69. Given [a scenario] determine which tree you would use, justify your answer. Examples:
 - You need to sort 1000 items.
 - You want to keep track of 10,000 key-value pairs such that you can (a) efficiently print all the items in key-order, (b) have fast “look-up”, and (c) relatively fast insertion.
 - You want to index a very large data set that does not fit into memory

Non-Balancing Search Trees (BST)

70. Why is the Big- O of inserting into a BST $O(n)$ and not $O(\log(n))$?
71. How can you determine the [min | max] value in a BST tree?
72. Given a node in a BST tree, how can you find [successor | predecessor] of that node (i.e. find the smallest value larger than the node or find the largest value smaller than the node).
73. Write the code for searching a BST for a given value. Assume a generic `Node<T>` class with a data field and left/right references. Assume the search method is given the root of the tree and a value to search for.
74. Given a BST and a value, show the steps to remove that value from the BST.
75. Write the code for [insert | remove | findMin | findMax | printSubset] in a BST. Assume a generic `Node<T>` class with a data field and `left/right` references.

Self-Balancing Search Trees (AVL / Red-Black)

76. Explain the “cases” for inserting into a [AVL | Red-Black] tree.
77. Determine an order of inserting the keys 1, 2, and 3 into an AVL which would require a [single right rotation | single left rotation | left-right double rotation | right-left double rotation].
78. Determine a set of keys and an order to insert them which would produce each of the cases for inserting into a Red-Black tree. Label each case with a meaningful name (not just “case 1”, “case 2”, etc.).
79. Compare AVL trees and Red-Black trees for pros and cons. Explain when/why you would use an AVL instead of a Red-Black Tree and explain when/why you would use a Red-Black Tree instead of an AVL.

Union Find / Disjoint Sets

80. Explain the types of problems the union-find data structure is designed to help with. Hint: why is it called “union-find”.
81. Given an array representing the results of a series of unions and finds, show the forest (graph form) it represents.
82. Explain the difference between “naive-union” and “rank union”.
83. Explain the difference between “naive-find” and “path compression find”.
84. Given a series of union and find operations, show the effects on the array which backs the union-find data structure using [naive | rank] union and [naive | path compression] find.
85. What is special about the performance analysis of union-find? Hint: this has to do with path compression and the definition of $\log^*(n)$.